

## Motivations

- Autoregressive generative models can estimate complex data distributions such as language, audio and images.
- Most existing methods either operate on discrete data distribution or discretize continuous data into several bins and use categorical distributions over the bins to approximate the continuous data distribution.
- Such approximation cannot express sharp changes in density without using significantly more bins, making it parameter inefficient.

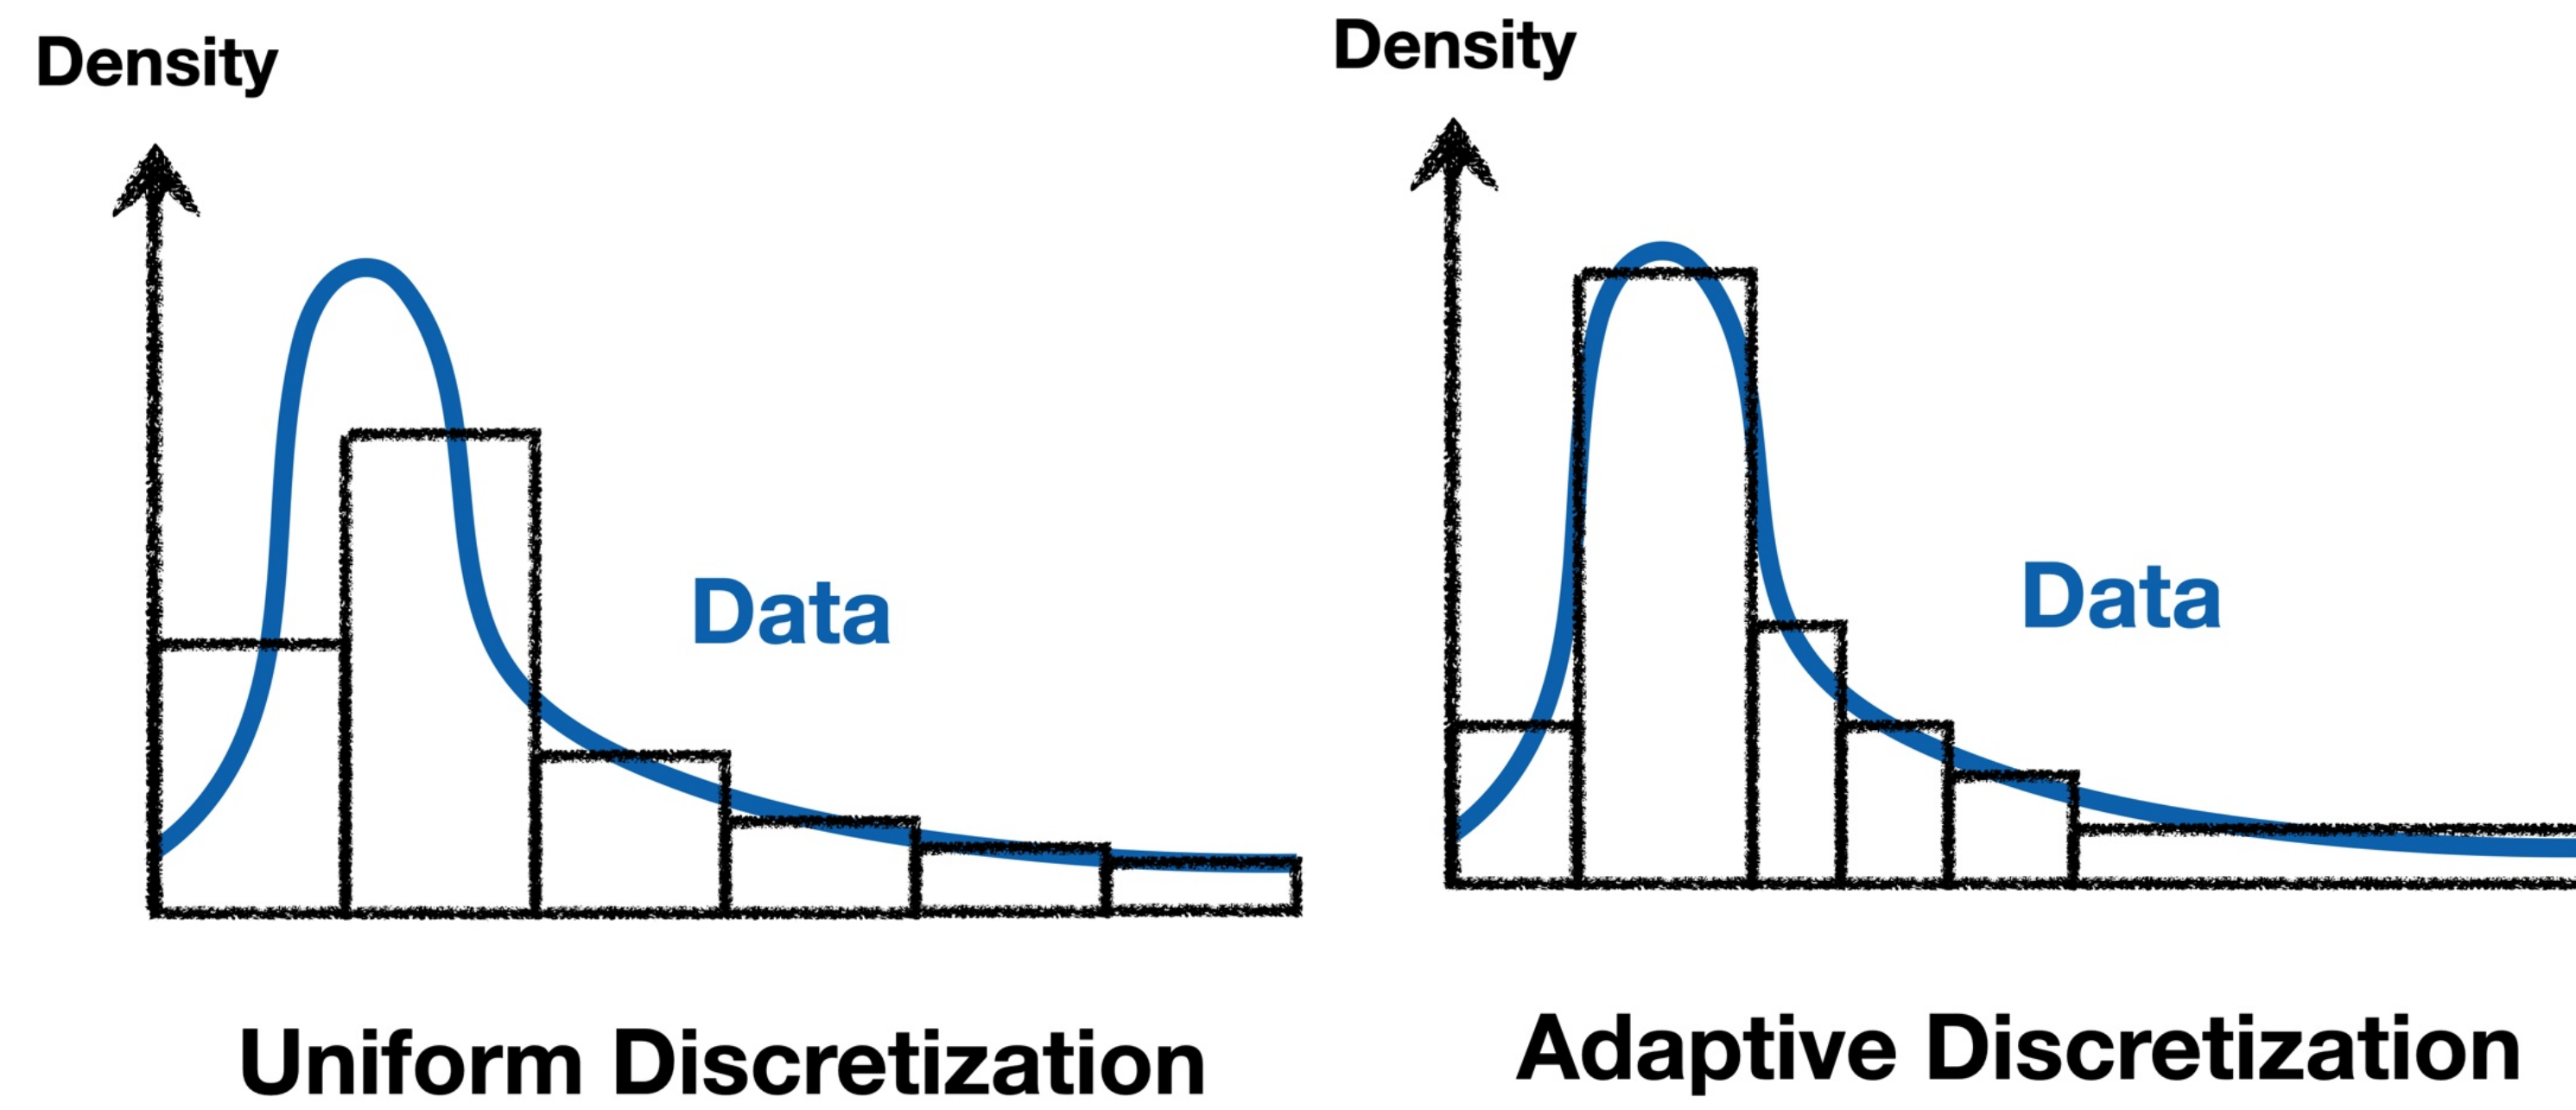
## Adaptive Discretization

AdaCat is a mixture of  $k$  non-overlapping truncated uniforms ( $w, h \in \mathbb{R}^k$ )

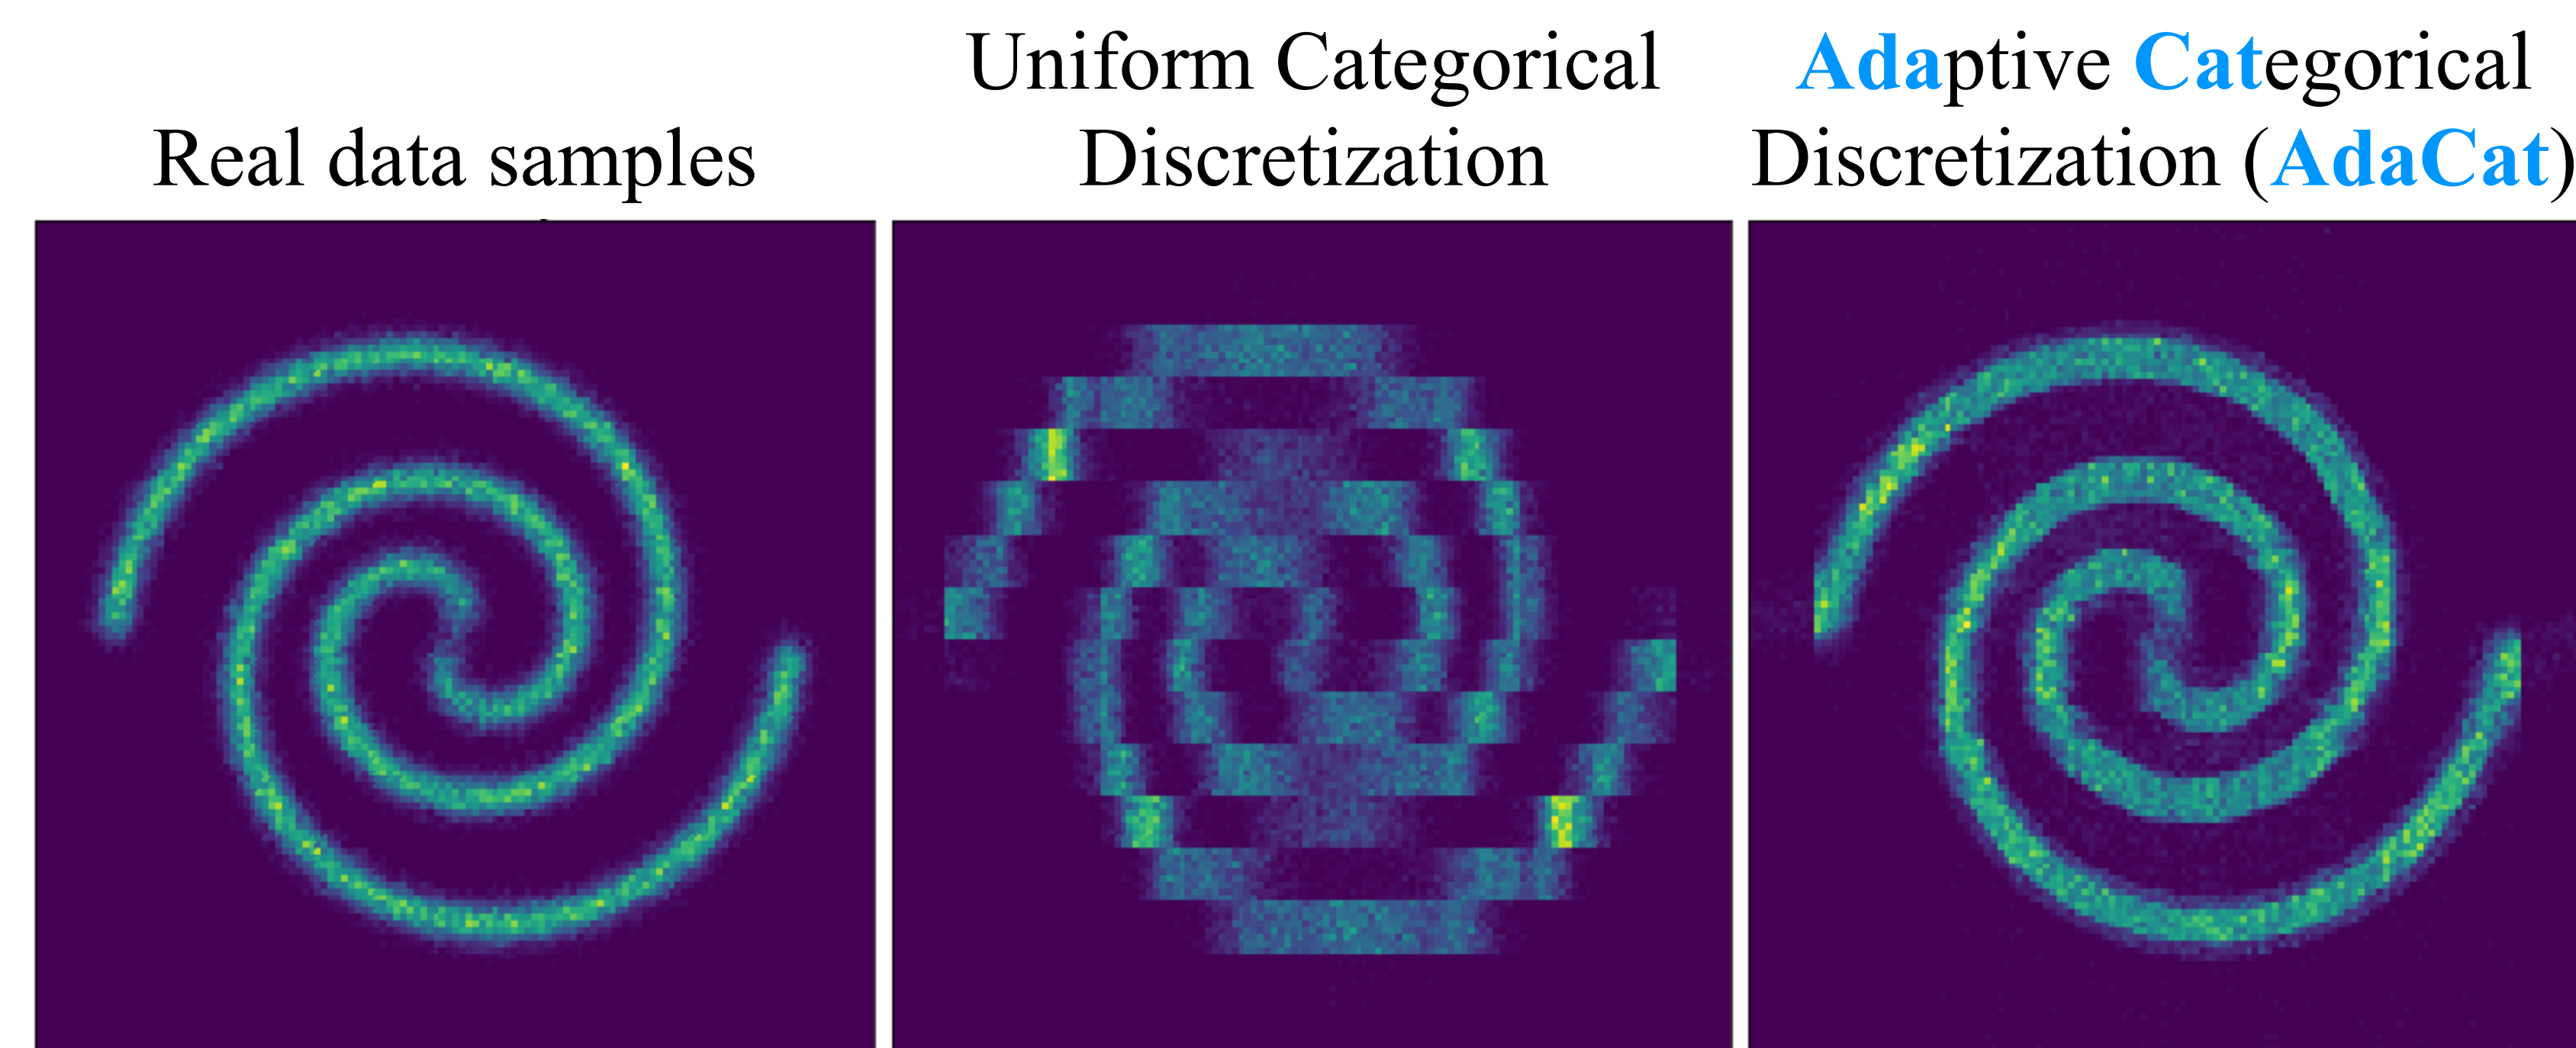
$$\text{AdaCat}_k(w, h) : f_{w,h,k}(x) = \sum_{i=1}^k \left\{ \mathbb{I}_{[c_i \leq x < c_i + w_i]} \frac{h_i}{w_i} \right\}, c_i = \sum_{j=1}^{i-1} w_j \quad (1)$$

- If  $w$  is fixed: Uniform discretization
- If  $h$  is fixed: Quantile regression

1D Illustrative Example ( $k = 6$ )



2D Toy Density Estimation ( $k = 12$ )



## Overcoming Non-Smoothness

### Empirical Negative Log-likelihood Objective

$$\hat{\mathcal{L}}_{\text{orig}}(\theta) = \frac{1}{n} \sum_{d=1}^n \sum_{t=1}^m \log p_{\theta}(x_d^t | x_d^{<t}) = \frac{1}{n} \sum_{d=1}^n \sum_{t=1}^m \log f_{w^t, h^t, k}(x_d^t)$$

where  $w^t, h^t = \text{nn}(x_d^{<t})$  are the predicted parameters of the AdaCat distribution.

- The gradient of  $\hat{\mathcal{L}}_{\text{orig}}$  with respect to the  $w^t, h^t$  is ill-defined at the bin boundaries.
- Therefore,  $\nabla_{\theta} \hat{\mathcal{L}}_{\text{orig}}$  might be biased. Optimizing  $\hat{\mathcal{L}}_{\text{orig}}$  can lead to bin collapse (red below).

### Smooth Objective

$$\hat{\mathcal{L}}_{\text{smooth}} = \frac{1}{n} \sum_{d=1}^n \sum_{t=1}^m \left[ \int_{\tilde{x}} \zeta(\tilde{x} | x_d^t) \log p_{\theta}(\tilde{x} | x_d^{<t}) d\tilde{x} \right]$$

where  $\zeta(\tilde{x} | x_d^t)$  is a smoothing distribution.


The inner integral can be computed as follows:

$$\int_{\tilde{x}} \zeta(\tilde{x}) \log f_{w,h,k}(\tilde{x}) d\tilde{x} = \sum_{j=1}^k [(F(c_j + w_j) - F(c_j))(\log h_j - \log w_j)]$$

where  $F(x)$  is the cumulative density function (CDF) of  $\xi(x)$ .

- The integral can be easily computed for  $\zeta$ 's when their CDFs are simple (uniform and Gaussian).
- The gradient is well-defined everywhere and  $\mathcal{L}_{\text{smooth}}$  prevents bin collapse (blue below).

### Optimization Dynamics of Original vs. Smooth Objective



## Empirical Results

### MNIST Density Estimation

| Parameters | Uniform      | Adaptive Quantile | DMoL [2]     | AdaCat       |
|------------|--------------|-------------------|--------------|--------------|
| 512        | N/A          | ×                 | 0.761        | <b>0.561</b> |
| 256        | <b>0.561</b> | ×                 | 0.698        | 0.573        |
| 216        | 0.838        | ×                 | 0.704        | <b>0.615</b> |
| 180        | 1.061        | ×                 | 0.684        | <b>0.629</b> |
| 152        | 1.299        | ×                 | 0.776        | <b>0.612</b> |
| 128        | 1.490        | ×                 | 0.700        | <b>0.608</b> |
| 64         | 2.453        | ×                 | 0.720        | <b>0.695</b> |
| 32         | 3.392        | 1.276             | <b>0.715</b> | 0.793        |
| Best       | <b>0.561</b> | 1.276             | 0.715        | <b>0.561</b> |

×: training diverged

### Offline RL with Model-based Planning in D4RL (Trajectory Transformer [1])

| Dataset            | Uniform                          | Quantile       | AdaCat                            |
|--------------------|----------------------------------|----------------|-----------------------------------|
| HalfCheetah-Medium | 44.0 $\pm 0.31$                  | 46.9 $\pm 0.4$ | <b>47.8 <math>\pm 0.22</math></b> |
| Hopper-Medium      | 67.4 $\pm 2.9$                   | 61.1 $\pm 3.6$ | <b>69.2 <math>\pm 4.5</math></b>  |
| Walker2d-Medium    | <b>81.3 <math>\pm 2.1</math></b> | 79.0 $\pm 2.8$ | 79.3 $\pm 0.8$                    |

See our paper for more results on tabular and audio data: [arxiv.org/abs/2208.02246](https://arxiv.org/abs/2208.02246)

## Try it out on PyTorch! – pip install adacat

```
from adacat.torch import Adacat
params = torch.nn.Parameter(torch.randn(10 * 2))
optim = torch.optim.Adam([params], lr=0.01)

for _ in range(n_its):
    dist = Adacat(params) # PyTorch distribution
    xs = sample_batch() # Sample from data distribution
    loss = -dist.log_prob(xs, smooth_coeff=0.001).mean()
    optim.zero_grad()
    loss.backward() # Optimize the smooth loss
    optim.step()
```

## References

- M. Janner, Q. Li, and S. Levine. Offline reinforcement learning as one big sequence modeling problem. In *Advances in Neural Information Processing Systems*, 2021.
- T. Salimans, A. Karpathy, X. Chen, and D. P. Kingma. Pixelcnn++: Improving the pixelcnn with discretized logistic mixture likelihood and other modifications. *arXiv preprint arXiv:1701.05517*, 2017.